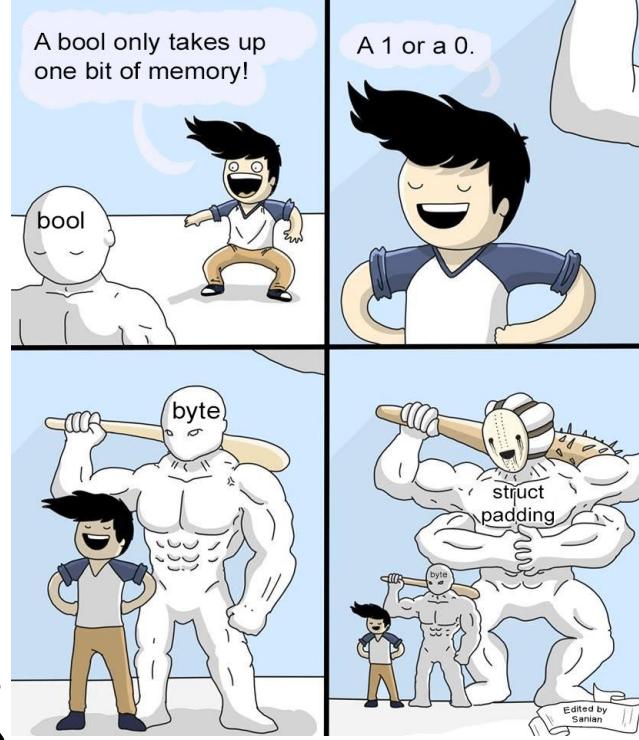


EECS 370 - Lecture 5

C to Assembly



Announcements

- Project 1
 - P1a due Thursday **1/29**
 - P1s due Thursday **2/5**
 - P1m due Thursday **2/5**
- HW1
 - Posted, due Monday **2/2**
- Lab
 - **Lab quiz due TODAY 1/22, 11:55pm**
 - Lab meets this upcoming Friday/Monday/Tuesday
- Midterm:
 - Thursday March 12th, 7-9pm
 - *Exam conflict form is due by January 23rd via the [Exam Conflict form](#) (ed post #17)*



Lecture 3 covered what you need.



Each problem is labelled by what lecture you need.

Agenda—last lecture

- ARM overview and basic instructions
- Memory instructions
 - **Handling multiple data widths**
- Sample Problems



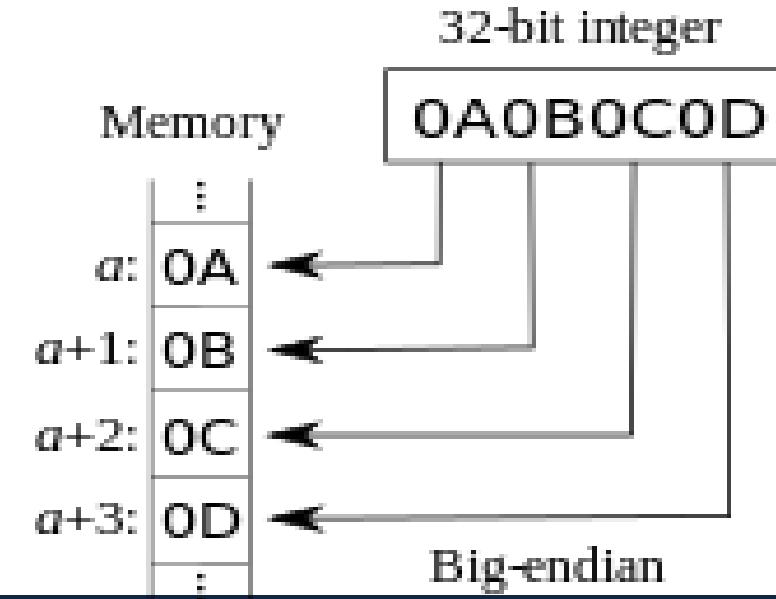
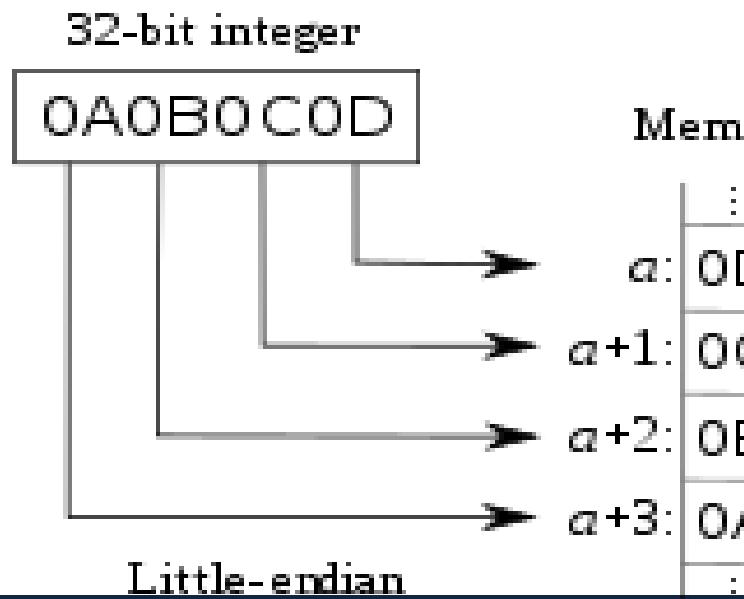
Load Instruction Sizes

How much data is retrieved from memory at the given address?

Desired amount of data to transfer?	Operation	Unused bits in register?	Example
64-bits (double word or whole register)	LDUR (Load unscaled to register)	N/A	0xFEDC_BA98_7654_3210
16-bits (half-word) into lower bits of reg	LDURH	Set to zero	0x0000_0000_0000_3210
8-bits (byte) into lower bits of reg	LDURB	Set to zero	0x0000_0000_0000_0010
32-bits (word) into lower bits of reg	LDURSW (load signed word)	Sign extend (0 or 1 based on most significant bit of transferred word)	0x0000_0000_7654_3210 or 0xFFFF_FFFF_F654_3210 (depends on bit 31)

Big Endian vs. Little Endian

- **Endian-ness:** ordering of bytes within a word
 - Little – Bigger address holds more significant bits
 - Big – Opposite, smaller address hold more significant bits
 - The Internet is big endian, x86 is little endian, LEG and ARMv8 can switch
 - But in general assume little endian. (Figures from Wikipedia)



Agenda—last lecture

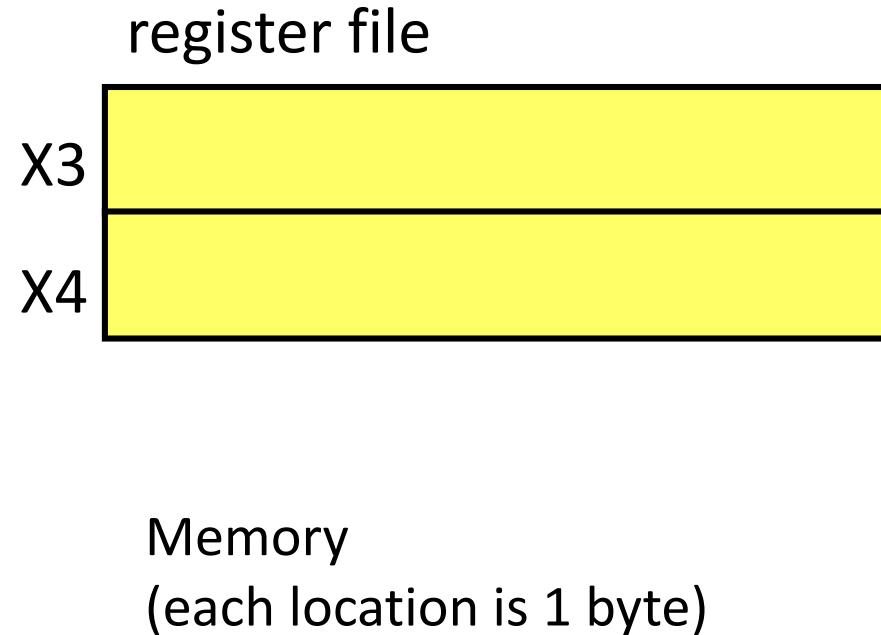
- ARM overview and basic instructions
- Memory instructions
 - Handling multiple data widths
- **Sample Problems**



Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

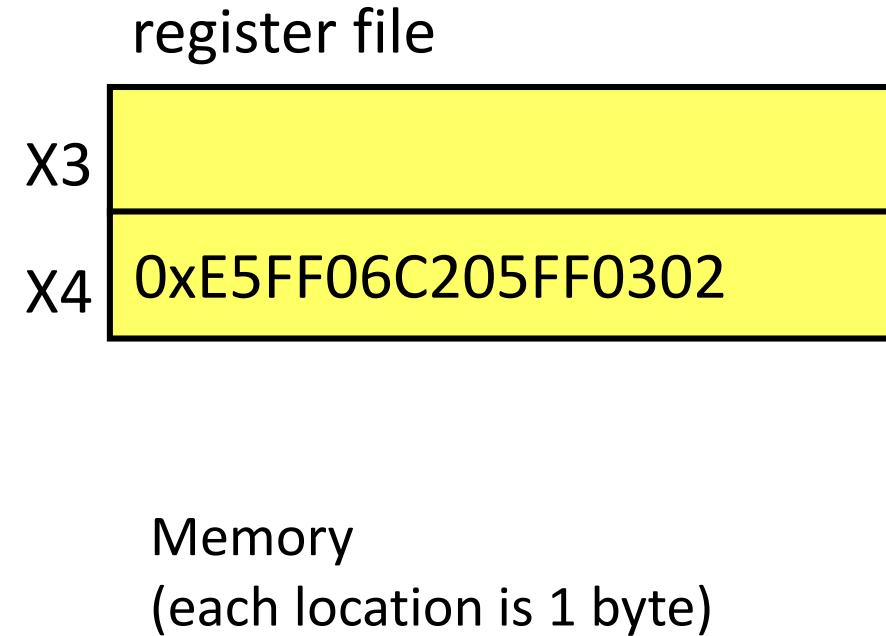
```
LDUR    X4, [X5, #100]
LDURB   X3, [X5, #102]
STUR    X3, [X5, #100]
STURB   X4, [X5, #102]
```



Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

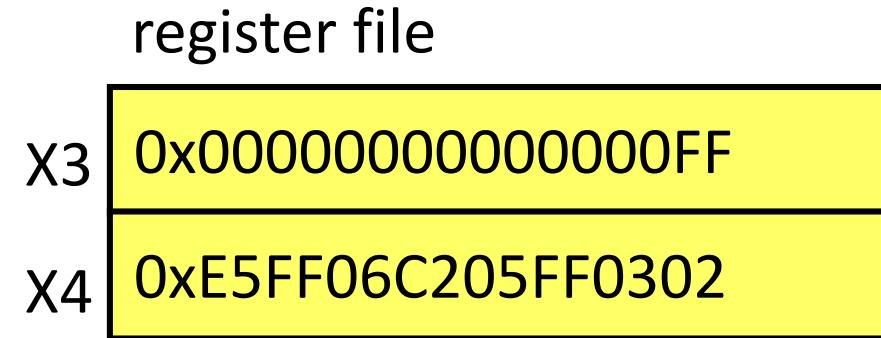
LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STUR	X3, [X5, #100]
STURB	X4, [X5, #102]



Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STUR	X3, [X5, #100]
STURB	X4, [X5, #102]



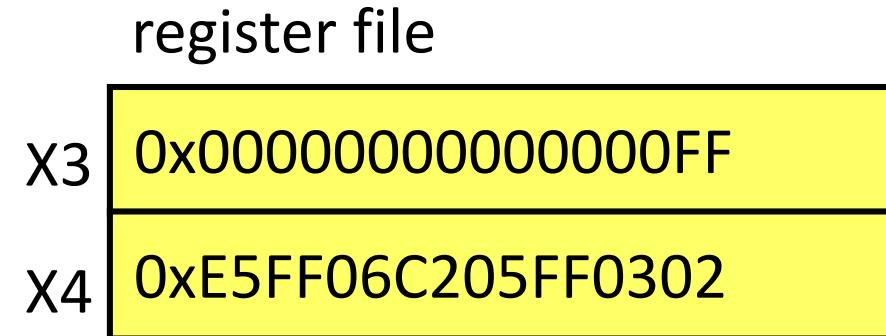
Memory
(each location is 1 byte)

little endian
0x02
0x03
0xFF
0x05
0xC2
0x06
0xFF
0xE5

Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STUR	X3, [X5, #100]
STURB	X4, [X5, #102]



Memory
(each location is 1 byte)

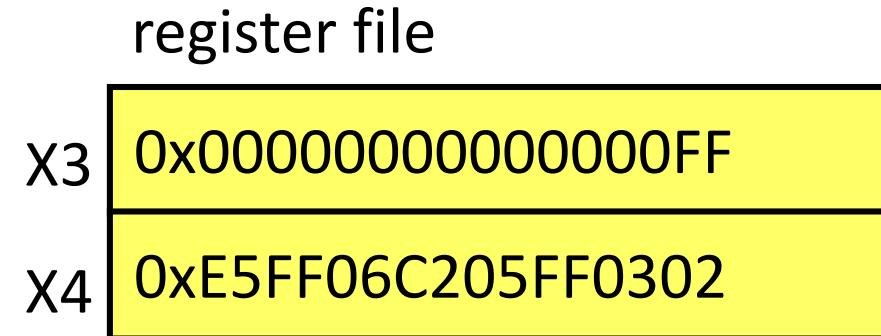
little endian

0xFF	100
0x00	101
0x00	102
0x00	103
0x00	104
0x00	105
0x00	106
0x00	107

Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STUR	X3, [X5, #100]
STURB	X4, [X5, #102]



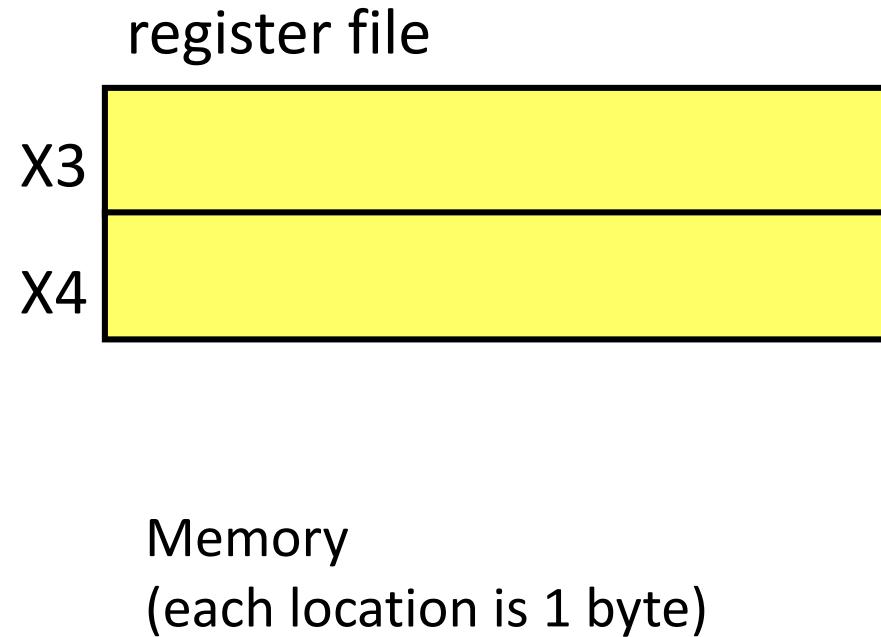
Memory
(each location is 1 byte)

little endian
0xFF
100
0x00
101
0x02
102
0x00
103
0x00
104
0x00
105
0x00
106
0x00
107

Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STURB	X3, [X5, #103]
LDURSW	X4, [X5, #100]

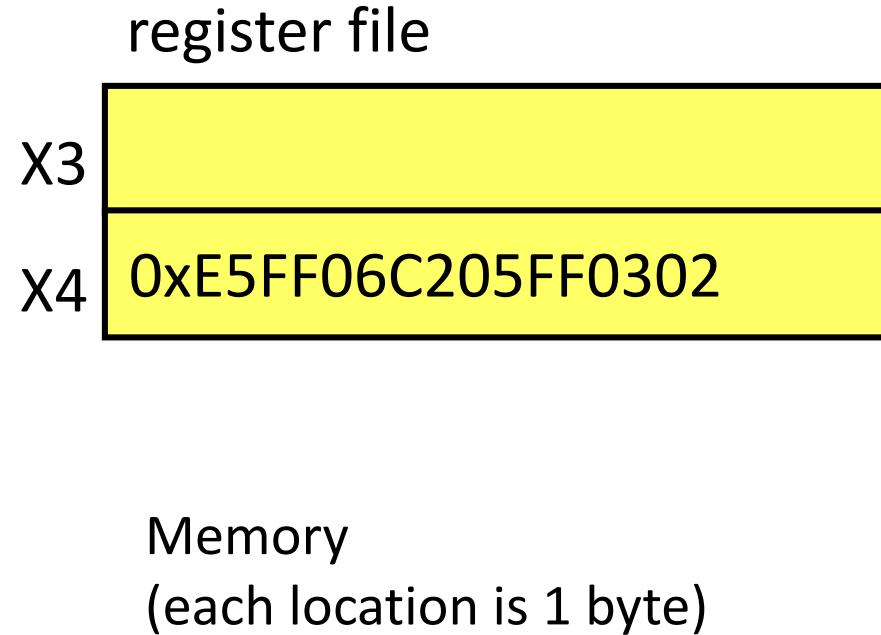


We show the registers as blank. What do they actually contain before we run the snippet of code?

Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STURB	X3, [X5, #103]
LDURSW	X4, [X5, #100]



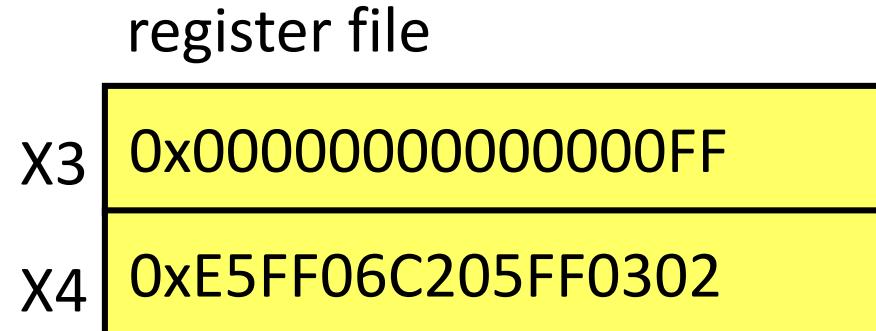
little endian

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STURB	X3, [X5, #103]
LDURSW	X4, [X5, #100]



Memory
(each location is 1 byte)

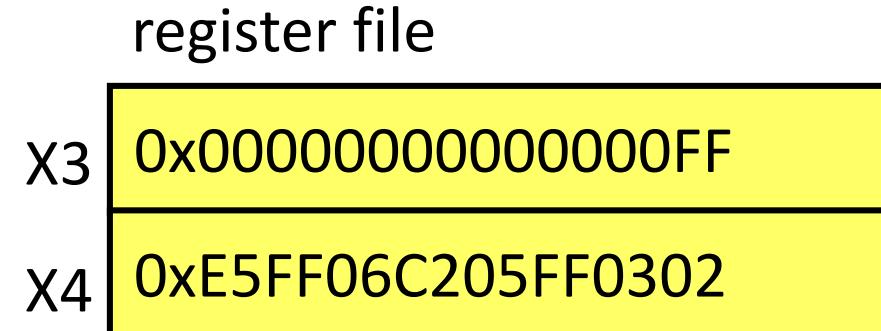
little endian

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STURB	X3, [X5, #103]
LDURSW	X4, [X5, #100]



Memory
(each location is 1 byte)

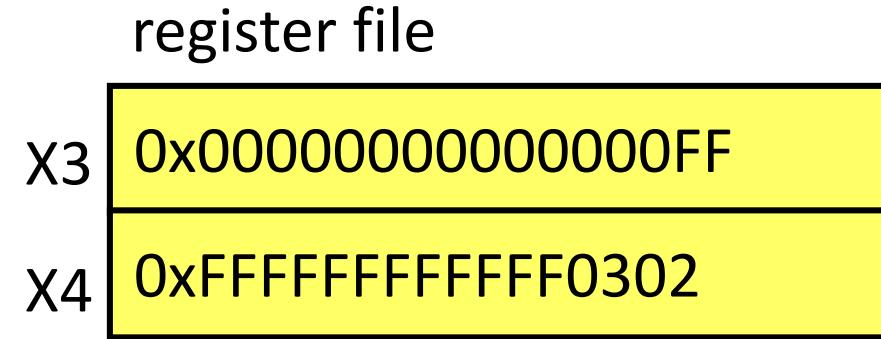
little endian

0x02	100
0x03	101
0xFF	102
0xFF	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

LDUR	X4, [X5, #100]
LDURB	X3, [X5, #102]
STURB	X3, [X5, #103]
LDURSW	X4, [X5, #100]



Memory
(each location is 1 byte)

little endian

0x02	100
0x03	101
0xFF	102
0xFF	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- **Lecture 5 : Converting C to assembly – basic blocks**
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout



Agenda

- **Memory alignment**
 - Aligning Structs
- Control flow instructions
 - C-code examples
- Extra Problems



Datatype	size (bytes)
char	1
short	2
int	4
double	8

Calculating Load/Store Addresses for Variables

```

short   a[100];
char    b;
int     c;
double  d;
short   e;
struct {
    char  f;
    int   g[1];
    char  h;
} i;

```

- *Problem:* Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed

$$a = 2 \text{ bytes} * 100 = 200$$

$$b = 1 \text{ byte}$$

$$c = 4 \text{ bytes}$$

$$d = 8 \text{ bytes}$$

$$e = 2 \text{ bytes}$$

$$f = 1 + 4 + 1 = 6 \text{ bytes}$$

total = 221, right or wrong?

Memory layout of variables

- Compilers don't like to place variables in memory arbitrarily
- As we'll see later in the course, memory is divided into fixed-sized **chunks**
 - When we load from a particular chunk, we really read the whole chunk
 - Usually an integer number of words (32 bits)
- If we read a single char (1 byte), it doesn't matter where it's placed

0x1000	0x1001	0x1002	0x1003
'a'	'b'	'c'	'd'

ldurb [x0, 0x1002]

- Reads [0x1000-0x1003], then throws away all but 0x1002, **fine**

Memory layout of variables

- BUT, if we read a 32-bit integer word, and that word starts at 0x1002:

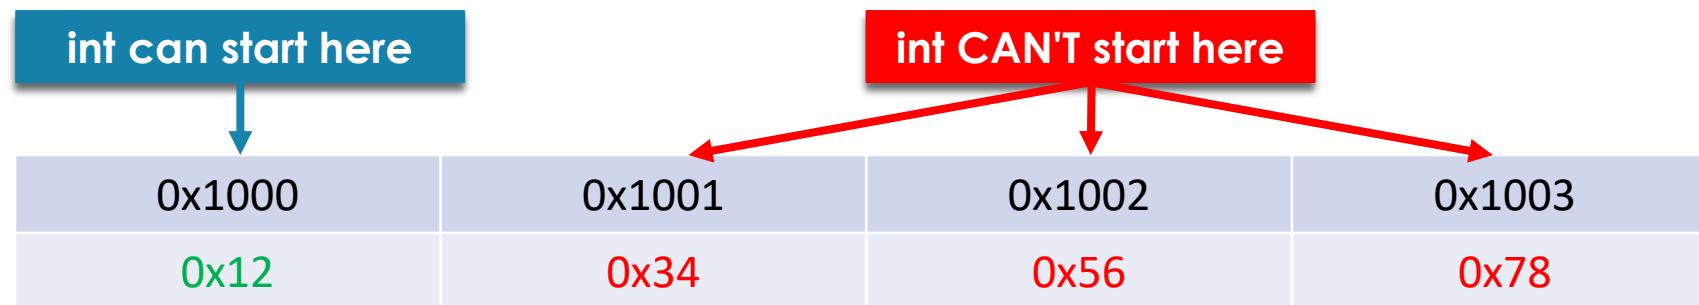
0x1000	0x1001	0x1002	0x1003
0xFF	0xFF	0x12	0x34
0x1004	0x1005	0x1006	0x1007
0x56	0x78	0xFF	0xFF

- First we need to read [0x1000-0x1003], throw away 0x1000 and 0x1001, **then** read [0x1004-0x1007]
- Need to read from memory twice! Slow! Complicated! **Bad!**

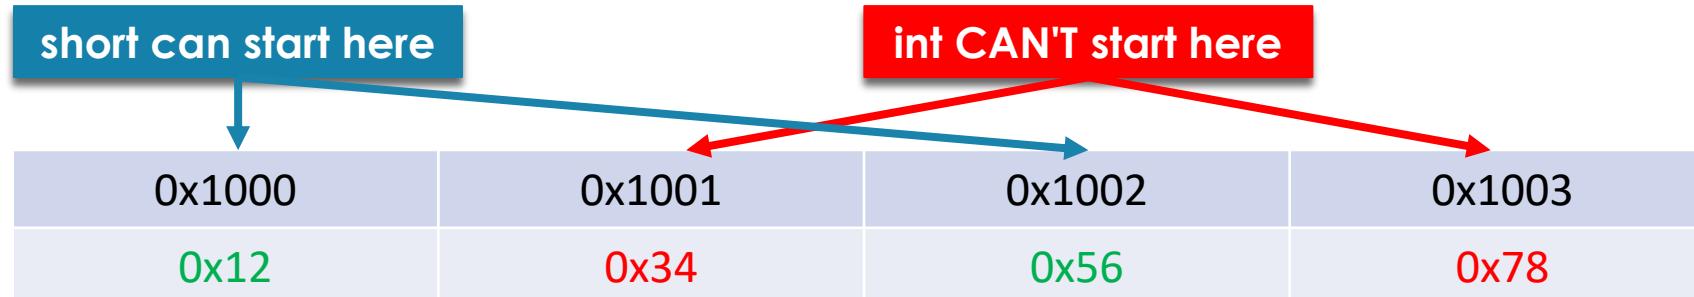
Solution: Memory Alignment

Where can chars start?

- Most modern ISAs require that data be aligned
 - An N-byte variable must start at an address A, such that $(A\%N) == 0$
- For example, starting address of a 32-bit **int** must be divisible by 4



- Starting address of a 16-bit **short** must be divisible by 2



Golden Rule of Alignment

- Every (primitive) object starts at an address divisible by its size
- "Padding" is placed in between objects if needed

```
char    c;  
short   s;  
int    i;
```

0x1000	0x1001	0x1002	0x1003	0x1004	0x1005	0x1006	0x1007
[c]	[padding]		[s]				[i]

- But what about non-primitive data types?
 - Arrays? Treat as independent objects
 - Structs? Trickier...

Agenda

- Memory alignment
 - Aligning Structs
- Control flow instructions
 - C-code examples
- Extra Problems



Problem with Structs

- If we align each element of a struct according to the Golden Rule, we can still run into issues
 - E.g.: An array of structs

Amount of padding
is different across
different instances

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	s[0].c	[pad]	[pad]	s[0].i				s[1].c	[pad]	[pad]	[pad]	s[1].i			

- Why is this bad?
- It makes "for" loops very difficult to write!
 - Offsets need to be different on each iteration

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

Structure Alignment

- Solution: in addition to laying out each field according to Golden Rule...
 - Identify largest (primitive) field
 - Starting address of overall struct is aligned based on the largest field
 - Padded in the back so total size is a multiple of the largest primitive

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	[pad]	[pad]	[pad]	s[0].c	[pad]	[pad]	[pad]		s[0].i		s[1].c	[pad]	[pad]	[pad]	

Guaranteed to lay
out each instance
identically



Structure Example

```
struct {  
    char w;  
    int x[3];  
    char y;  
    short z;  
}
```

What boundary should this struct be aligned to?

- a) 1 byte
- b) 4 bytes
- c) 12 bytes
- d) 2 bytes
- e) 19 bytes

- Assume struct starts at location 1000,
 - char w → 1000
 - x[0] → 1004–1007, x[1] → 1008 – 1011, x[2] → 1012 – 1015
 - char y → 1016
 - short z → 1018 – 1019

Total size = 20 bytes!

Datatype	size (bytes)
char	1
short	2
int	4
double	8

Calculating Load/Store Addresses for Variables

```

short a[100];
char b;
int c;
double d;
short e;
struct {
    char f;
    int g[1];
    char h;
} i;

```

- *Problem:* Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed
 - a = 200 bytes (100-299)
 - b = 1 byte (300-300)
 - c = 4 bytes (304-307)
 - d = 8 bytes (312-319)
 - e = 2 bytes (320-321)
 - struct: largest field is 4 bytes, start at 324
 - f = 1 byte (324-324)
 - g = 4 bytes (328-331)
 - h = 1 byte (332-332)
 - i = 12 bytes (324-335)
- 236 bytes total!! (compared to 221, originally)

Data Layout – Why?

- Does gcc (or another compiler) reorder variables in memory to avoid padding?
- Only outside structs
 - C99 forbids reordering elements inside a struct!
- The programmer (i.e., you) are expected to manage data layout of variables for your program and structs.
- Two optimal strategies:
 - Order fields in struct by datatype size, smallest first
 - Or by largest first



Agenda

- Memory alignment
 - Aligning Structs
- **Control flow instructions**
 - C-code examples
 - Extra Problems

ARM/LEGv8 Sequencing Instructions

- Sequencing instructions change the flow of instructions that are executed
 - This is achieved by modifying the program counter (PC)
- Unconditional branches are the most straightforward they ALWAYS change the PC and thus “jump” to another instruction out of the usual sequence
- Conditional branches

If (*condition_test*) goto *target_address*

condition_test examines the four flags from the processor status word (SPSR)

target_address is a 19 bit signed word displacement on current PC

LEGv8 Conditional Instructions

- Two varieties of conditional branches
 1. One type compares a register to see if it is equal to zero.
 2. Another type checks the condition codes set in the status register.

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch	
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch	
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	test condition codes; if true, branch	

- Let's look at the first type: CBZ and CBNZ
 - CBZ: Conditional Branch if Zero
 - CBNZ: Conditional Branch if Not Zero

LEGv8 Conditional Instructions

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
 - The relative address is a 19 bit signed integer—the number of instructions. Recall instructions are 32 bits of 4 bytes

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Example: CBNZ X3, Again
 - If X3 doesn't equal 0, then branch to label "Again"
 - "Again" is an offset from the PC of the current instruction (CBNZ)
 - Why does "25" in the above table result in PC + 100?

LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

Again: ADDI X3, X3, #-1

CBNZ X3, Again

What is the offset?

- a) -16
- b) -12
- c) -4
- d) -3
- e) 0

LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

Again: ADDI X3, X3, #-1

CBNZ X3, Again

- Answer = -3
 - The offset field is 19 bits signed so the bit pattern would be
111 1111 1111 1111 1101
 - Two 00's are appended to the above 19 bits and then the result would be sign-extended (with one's) to 64 bits and added to the value of PC at CBNZ
 - Why the two 00's?



LEGv8 Conditional Instructions

- Motivation:
 - Some types of branches makes sense to check if a certain value is zero or not
 - `while(a)`
 - But not all:
 - `if(a > b)`
 - `if(a == b)`
 - Using an extra **program status register** to check for various conditions allows for a greater breadth of branching behavior



LEGv8 Conditional Instructions Using FLAGS

- FLAGS: NZVC record the results of (arithmetic) operations Negative, Zero, oVerflow, Carry—not present in LC2K
- We explicitly set them using the “set” modification to ADD/SUB etc.
- Example: ADDS causes the 4 flag bits to be set according as the outcome is negative, zero, overflows, or generates a carry

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	$X1 = X2 + X3$	Three register operands
	subtract	SUB X1, X2, X3	$X1 = X2 - X3$	Three register operands
	add immediate	ADDI X1, X2, 20	$X1 = X2 + 20$	Used to add constants
	subtract immediate	SUBI X1, X2, 20	$X1 = X2 - 20$	Used to subtract constants
	add and set flags	ADDS X1, X2, X3	$X1 = X2 + X3$	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	$X1 = X2 - X3$	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	$X1 = X2 + 20$	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	$X1 = X2 - 20$	Subtract constant, set condition codes



ARM Condition Codes Determine Direction of Branch

- In LEGv8 only ADDS / SUBS / ADDIS / SUBIS / CMP / CMPI set the condition codes FLAGS or condition codes in PSR—the program status register
- Four primary condition codes evaluated:
 - N – set if the result is **negative** (i.e., bit 63 is non-zero)
 - Z – set if the result is **zero** (i.e., all 64 bits are zero)
 - C – set if last addition/subtraction had a **carry**/borrow out of bit 63
 - V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)
- Don't worry about the C and V for this class



ARM Condition Codes Determine Direction of Branch--continued

Encoding	Name (& alias)	Meaning (integer)	Flags
0000	EQ	Equal	$Z==1$
0001	NE	Not equal	$Z==0$
0010	HS (CS)	Unsigned higher or same (Carry set)	$C==1$
0011	LO (CC)	Unsigned lower (Carry clear)	$C==0$
0100	MI	Minus (negative)	$N==1$
0101	PL	Plus (positive or zero)	$N==0$
0110	VS	Overflow set	$V==1$
0111	VC	Overflow clear	$V==0$
1000	HI	Unsigned higher	$C==1 \&& Z==0$
1001	LS	Unsigned lower or same	$! (C==1 \&& Z==0)$
1010	GE	Signed greater than or equal	$N==V$
1011	LT	Signed less than	$N \neq V$
1100	GT	Signed greater than	$Z==0 \&& N==V$
1101	LE	Signed less than or equal	$! (Z==0 \&& N==V)$
1110	AL	Always	Any
1111	NV [†]		

Need to know the 7 with the red arrows

```
CMP X1, X2  
B.LE Label1
```

For this example,
we branch if X1 is
 \leq to X2

Conditional Branches: How to use

- CMP instruction lets you compare two registers.
 - Could also use SUBS etc.
 - That could save you an instruction.
- B.cond lets you branch based on that comparison.
- Example:

```
CMP    X1,  X2
B.GT  Label1
```

- Branches to Label1 if X1 is greater than X2.

Agenda

- Memory alignment
 - Aligning Structs
- Control flow instructions
 - **C-code examples**
- Extra Problems



Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

Using Labels		Without Labels
CMP X1, X2 B.NE L1 ADD X1, X1, #1 B L2 L1: ADD X2, X2, #1 L2: ...		CMP X1, X2 B.NE 3 ADD X1, X1, #1 B 2 ADD X2, X2, #1

Note that conditions in assembly are often the inverse of the "if" condition. Why?

Assemblers must deal with labels and assign displacements

Loop—Example

```
// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2
```

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

of branch instructions
= $3 \times 10 + 1 = 31$

a.k.a. while-do template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	CMPI	X1, #10
	B.EQ	endLoop
	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endif:	ADDI	X1, X1, #1
	B	Loop1
endLoop:		

Agenda

- Memory alignment
 - Aligning Structs
- Control flow instructions
 - C-code examples
- **Extra Problems**



Extra Example: Do-while Loop

```
// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2
```

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

of branch instructions
= $2 * 10 = 20$

a.k.a. do-while template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endIf:	ADDI	X1, X1, #1
	CMPI	X1, #10
	B.LT	Loop1
endLoop:		

Extra Example: Do-while Loop

```
// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2
```

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

of branch instructions

$$= 2 * 10 = 20$$

a.k.a. do-while template

Extra Problem – For Your Reference

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in x1
// struct {int val; struct node *next;} node;
// struct node *ptr;

if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```

Extra Problem

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in x1
// struct {int val; struct node *next;} node;
// struct node *ptr;

if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```

```
        cmp r1, #0
        beq Endif
        ldursw r2, [r1, #0]
        cmp r2, #0
        b.le Endif
        add r2, r2, #1
        str r2, [r1, #0]
Endif : ....
```

Extra Class Problem

- How much memory is required for the following data, assuming that the data starts at address 200 and is a 32 bit address space?

```
int a;  
struct {double b, char c, int d} e;  
char* f;  
short g[20];
```

How much memory?

- a) $x < 40$ bytes
- b) $40 < x < 50$ bytes
- c) $50 < x < 60$ bytes
- d) $60 < x$ bytes



Next Time

- More C-to-Assembly
 - Function calls

