

# EECS 370 - Lecture 6

## Function Calls



# Announcements

- Project 1
  - P1a due Thursday **1/29**
  - P1s due Thursday **2/5**
  - P1m due Thursday **2/5**
- HW1
  - Posted, due Monday **2/2**
- Lab
  - **Lab quiz due TODAY 1/22, 11:55pm**
  - Lab meets this upcoming Friday/Monday/Tuesday
- Midterm:
  - Thursday March 12<sup>th</sup>, 7-9pm
  - *Exam conflict form is due by January 23<sup>rd</sup> via the [Exam Conflict form](#) (ed post #17)*



**Lecture 3 covered what you need.**



**Each problem is labelled by what lecture you need.**

# Review of last lecture

- Memory alignment
  - **Aligning Structs**
- Control flow instructions
  - C-code examples



Datatype	size (bytes)
char	1
short	2
int	4
double	8

# Calculating Load/Store Addresses for Variables

```

short a[100];
char b;
int c;
double d;
short e;
struct {
    char f;
    int g[1];
    char h;
} i;

```

- *Problem:* Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed

a = 200 bytes (100-299)

b = 1 byte (300-300)

c = 4 bytes (304-307)

d = 8 bytes (312-319)

e = 2 bytes (320-321)

struct: largest field is 4 bytes, start at 324

f = 1 byte (324-324)

g = 4 bytes (328-331)

h = 1 byte (332-332)

i = 12 bytes (324-335)

236 bytes total!! (compared to 221, originally)

# Data Layout – Why?

- Does gcc (or another compiler) reorder variables in memory to avoid padding?
- Only outside structs
  - C99 forbids reordering elements inside a struct
- The programmer (i.e., you) are expected to manage data layout of variables for your program and structs.
- Two optimal strategies:
  - Order fields in struct by datatype size, smallest first
  - Or by largest first



# Agenda

- Memory alignment
  - Aligning Structs
- **Control flow instructions**
  - C-code examples



# LEGv8 Conditional Instructions

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
  - The relative address is a 19 bit signed integer—the number of instructions. Recall instructions are 32 bits of 4 bytes

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Example: CBNZ X3, Again
  - If X3 doesn't equal 0, then branch to label "Again"
  - "Again" is an offset from the PC of the current instruction (CBNZ)
  - Why does "25" in the above table result in PC + 100?

# LEGv8 Conditional Instructions

- Motivation:
  - Some types of branches makes sense to check if a certain value is zero or not
    - `while(a)`
  - But not all:
    - `if(a > b)`
    - `if(a == b)`
  - Using an extra **program status register** to check for various conditions allows for a greater breadth of branching behavior



# LEGv8 Conditional Instructions Using FLAGS

- FLAGS: NZVC record the results of (arithmetic) operations Negative, Zero, oVerflow, Carry—not present in LC2K
- We explicitly set them using the “set” modification to ADD/SUB etc.
- Example: ADDS causes the 4 flag bits to be set according as the outcome is negative, zero, overflows, or generates a carry

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	$X1 = X2 + X3$	Three register operands
	subtract	SUB X1, X2, X3	$X1 = X2 - X3$	Three register operands
	add immediate	ADDI X1, X2, 20	$X1 = X2 + 20$	Used to add constants
	subtract immediate	SUBI X1, X2, 20	$X1 = X2 - 20$	Used to subtract constants
	add and set flags	ADDS X1, X2, X3	$X1 = X2 + X3$	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	$X1 = X2 - X3$	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	$X1 = X2 + 20$	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	$X1 = X2 - 20$	Subtract constant, set condition codes



# ARM Condition Codes Determine Direction of Branch

- In LEGv8 only ADDS / SUBS / ADDIS / SUBIS / CMP / CMPI set the condition codes FLAGS or condition codes in PSR—the program status register
- Four primary condition codes evaluated:
  - N – set if the result is **negative** (i.e., bit 63 is non-zero)
  - Z – set if the result is **zero** (i.e., all 64 bits are zero)
  - C – set if last addition/subtraction had a **carry**/borrow out of bit 63
  - V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)
- Don't worry about the C and V for this class



# ARM Condition Codes Determine Direction of Branch--continued

Encoding	Name (& alias)	Meaning (integer)	Flags
0000	EQ	Equal	$Z==1$
0001	NE	Not equal	$Z==0$
0010	HS (CS)	Unsigned higher or same (Carry set)	$C==1$
0011	LO (CC)	Unsigned lower (Carry clear)	$C==0$
0100	MI	Minus (negative)	$N==1$
0101	PL	Plus (positive or zero)	$N==0$
0110	VS	Overflow set	$V==1$
0111	VC	Overflow clear	$V==0$
1000	HI	Unsigned higher	$C==1 \&& Z==0$
1001	LS	Unsigned lower or same	$! (C==1 \&& Z==0)$
1010	GE	Signed greater than or equal	$N==V$
1011	LT	Signed less than	$N \neq V$
1100	GT	Signed greater than	$Z==0 \&& N==V$
1101	LE	Signed less than or equal	$! (Z==0 \&& N==V)$
1110	AL	Always	Any
1111	NV <sup>†</sup>		

Need to know the 7 with the red arrows

```
CMP X1, X2  
B.LE Label1
```

For this example,  
we branch if X1 is  
 $\leq$  to X2

# Conditional Branches: How to use

- CMP instruction lets you compare two registers.
  - Could also use SUBS etc.
    - That could save you an instruction.
  - B.cond lets you branch based on that comparison.
- Example:

```
CMP    X1,  X2
B.GT  Label1
```

- Branches to Label1 if X1 is greater than X2.

# Agenda

- Memory alignment
  - Aligning Structs
- Control flow instructions
  - **C-code examples**
- Extra Problems



# Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

# Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

Using Labels		Without Labels
CMP X1, X2 B.NE L1 ADD X1, X1, #1 B L2 L1: ADD X2, X2, #1 L2: ...		CMP X1, X2 B.NE 3 ADD X1, X1, #1 B 2 ADD X2, X2, #1

Note that conditions in assembly are often the inverse of the "if" condition. Why?

Assemblers must deal with labels and assign displacements

# Loop—Example

```
// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2
```

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

# of branch instructions  
=  $3 \times 10 + 1 = 31$

a.k.a. while-do template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	CMPI	X1, #10
	B.EQ	endLoop
	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endif:	ADDI	X1, X1, #1
	B	Loop1
endLoop:		

Note: Could further optimize by counting down (9 to 0) and using CNBZ in place of the CMPI and B.LT at the end!

# Extra Example: Do-while Loop

```
// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2
```

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

# of branch instructions  
=  $2 * 10 = 20$

a.k.a. do-while template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endIf:	ADDI	X1, X1, #1
	CMPI	X1, #10
	B.LT	Loop1
endLoop:		

# Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- **Lecture 6 : Converting C to assembly – functions**
- Lecture 7 : Translation software; libraries, memory layout



# Agenda

- **Branching far away**
- Function calls and the call stack
- Assigning variables to memory locations
- Saving registers
- Caller/callee example

# Branching far away

- Underlying philosophy of ISA design: **make the common case fast**
- Most branches target nearby instructions
  - Displacement of 19 bits is usually enough
- BUT what if we need to branch really far away (more than  $2^{19}$  words)?  
`CBZ X15, FarLabel`
- The assembler is smart enough to replace that with  
`CBNZ X15, L1`  
`B FarLabel`  
`L1: .....`
- The simple branch instruction (B) has a 26 bit offset which spans about 64 million instructions!
- In LC2K, we can do a similar thing by using JALR instead of BEQ

# Agenda

- Branching far away
- **Function calls and the call stack**
- Assigning variables to memory locations
- Saving registers
- Caller/callee example

# Implementing Functions

**Poll: What's wrong  
with this?**

- Does this assembly code do what we need?

```
int mult_2(int x){  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main(){  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

```
LDURSW X1, [XZR, GLOBAL]  
ADD    X2, X1, #1           // Inc GLOBAL  
STURW X2, [XZR, X]         // Pass arg  
B      MULT_2              // Execute func  
  
RETURN:  
LDURSW X3, [XZR, RETURN]   // load result  
STURW X3, [XZR, STRING]   // Pass arg  
B PRINTF                  // Execute func  
...  
MULT_2:  
LDURSW X1, [XZR, X]         // load arg  
ADD    X2, X1, X1           // mult by 2  
STURW X2, [XZR, TEMP]      // return result  
B RETURN                   // return
```

# Problem 1: Returning from Functions

- Branches so far have hard-coded destination

```
B.NE L1  
ADD X1, X1, #1  
B L2  
L1: ADD X2, X2, #1  
L2: ...
```

```
B.NE      3  
ADD      X1, X1, #1  
B        2  
ADD      X2, X2, #1
```

```
int func(int x) {  
    printf(x * 10);  
    return;  
}  
int helper() {  
    func(7);  
}  
int main() {  
    helper();  
    func(13);  
}
```

Should this return to "helper" or "main"?

- This is fine for if-statements, for-loops etc
- But functions can be called from multiple places
  - Meaning we'll return to different spots on each func call! Can't hardcode offset!

# Solution: Indirect Jumps

- Indirect branches or "jumps" don't hardcode destination in register
- They index a register whose value holds destination

Unconditional branch	branch	B 2500	go to PC + 10000	Branch to target address; PC-relative
	branch to register	BR X30	go to X30	For switch, procedure return
	branch with link	BL 2500	X30 = PC + 4; PC + 10000	For procedure call PC-relative

- Use "**BL**" to **call a function**
  - Destination is hardcoded
  - PC +4 (return address) stored in X30
- Use "**BR**" to **return from a function**
  - X30 is read for return address
  - Allows us to return to different places



# Solution: Indirect Jumps

```
int mult_2(int x){  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main(){  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

Also don't  
need "return"  
labels

```
LDURSW X1, [XZR, GLOBAL]  
ADD X2, X1, #1 // Inc GLOBAL  
STURW X2, [XZR, X] // Pass arg  
BL MULT_2 // Execute func  
  
RETURN:  
LDURSW X3, [XZR, RETURN] // load result  
STURW X3, [XZR, STRING] // Pass arg  
BL PRINTF // Execute func  
...  
MULT_2:  
LDURSW X1, [XZR, X] // load arg  
ADD X2, X1, X1 // mult by 2  
STURW X2, [XZR, TEMP] // return result  
BR // return
```

Now MULT\_2  
can return to  
whatever  
function  
called it

# Problem 2: Passing Parameters

For any recursive functions, global variables will be overwritten

```
int mult_2(int x){  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
...  
  
int main(){  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

```
LDURSW X1, [XZR, GLOBAL] // Load GLOBAL  
ADD X2, X1, #1           // Inc GLOBAL  
STURW X2, [XZR, X]       // Pass arg  
BL MULT_2                // Execute func  
LDURSW X3, [XZR, RETURN] // Load result  
STURW X3, [XZR, STRING]  // Pass arg  
BL PRINTF                 // Execute func  
  
...  
  
MULT_2:  
LDURSW X1, [XZR, X]       // Load arg  
ADD X2, X1, X1            // mult by 2  
STURW X2, [XZR, TEMP]     // Return result  
BR                         // Return
```

How to pass in parameters?



# Task 1: Passing parameters

- Where should you put all of the parameters?
  - Registers?
    - Fast access but few in number and wrong size for some objects
  - Memory?
    - Good general solution but slow
- ARMv8 solution—and the usual answer:
  - Both
    - Put the first few parameters in registers (if they fit) (X0 – X7)
    - Put the rest in memory on the call stack— **important concept**

# Call stack

- ARM conventions (and most other processors) allocate a region of memory for the “call” stack
  - This memory is used to manage all the storage requirements to simulate function call semantics
    - Parameters (that were not passed through registers)
    - Local variables
    - Temporary storage (when you run out of registers and need somewhere to save a value)
    - Return address
    - Etc.
- Sections of memory on the call stack [**stack frames**] are allocated when you make a function call, and de-allocated when you return from a function

# The stack grows as functions are called

FUNCTION CALLS

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

inside foo

foo's stack frame

bar calls printf

foo's stack frame

bar's stack frame

printf's stack frame

foo calls bar

foo's stack frame

bar's stack frame

# The stack shrinks as functions return

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

printf returns

foo's stack frame

bar's stack frame

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

bar returns

foo's stack frame

# Stack frame contents

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}  
  
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

foo's stack frame


return addr to main
x
y[0]
y[1]
spilled registers in foo

# Stack frame contents (2)

```
void foo()
{
    int x, y[2];
    bar(x);
}

void bar(int x)
{
    int a[3];
    printf();
}
```

**Spill data**—not enough room in x0-x7 for params and also caller and callee saves



# Agenda


- Branching far away
- Function calls and the call stack
- **Assigning variables to memory locations**
- Saving registers
- Caller/callee example

## Review: Where do the variables go?

# Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```




# Assigning variables to memory spaces

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```


w goes in static, as it's a global  
x goes on the stack, as it's a parameter  
  
y goes in static, 1 copy of this!!  
p goes on the stack  
allocate 10 bytes on heap, ptr  
set to the address  
string literal "%s\n" goes in static,  
implicit pointer to string on stack, p goes  
on stack

The addresses of local variables  
will be different depending on  
where we are in the call stack



# Accessing Local Variables

- Stack pointer (SP):
  - register that keeps track of current top of stack
- Compiler (or assembly writer) knows relative offsets of objects in stack
- Can access using lw/sw offsets



# Agenda

- Branching far away
- Function calls and the call stack
- Assigning variables to memory locations
- **Saving registers**
- Caller/callee example

# Problem 3: Reusing registers

- Higher level languages (like C/C++) provide many abstractions that don't exist at the assembly level
- E.g. in C, each function has its own local variables
  - Even if different functions have local variables with the same name, they are independent and guaranteed not to interfere with each other!


```
void foo() {  
    int a=1;  
    bar();  
    printf(a);  
}
```

Still prints "1"...  
these don't  
interfere

```
void bar() {  
    int a=2;  
    return;  
}
```

# What about registers?

- But in assembly, all functions share a small set (e.g. 32) of registers
  - Called functions will overwrite registers needed by calling functions



- "Someone" needs to save/restore values when a function is called to ensure this doesn't happen

# Two Possible Solutions

- Either the **called** function **saves** register values before it overwrites them and **restores** them before the function returns (**callee saved**)...

```
main: movz x0, #1
      bl foo
      bl printf
```

```
foo: stur x0, [stack]
      movz x0, #2
      ldur x0, [stack]
      br x30
```

- Or the **calling** function **saves** register values before the function call and **restores** them after the function call (**caller saved**)...

```
main: movz x0, #1
      stur x0, [stack]
      bl foo
      ldur x0, [stack]
      bl printf
```

```
foo: movz x0, #2
      br x30
```

# Another example

Original C  
Code

```
void foo(){  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

No need to  
save r2/r3.  
Why?

Additions for  
Caller-save

```
void foo(){  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
save r1 to stack  
save r4 to stack  
bar();  
restore r4  
restore r1  
d = a+d;  
return();  
}
```

Assume bar() will  
overwrite registers  
holding a,d

Additions for  
Callee-save

```
void foo(){  
    int a,b,c,d;  
save r1  
save r2  
save r3  
save r4  
a = 5; b = 6;  
c = a+1; d=c-1;  
bar();  
d = a+d;  
restore r4  
restore r3  
restore r2  
restore r1  
return();  
}
```

bar() will save a,b, but  
now foo() must save  
main's variables

# “caller-save” vs. “callee-save”

- Caller-save
  - What if bar() doesn't use r1/r4?
  - No harm done, but wasted work
- Callee-save
  - What if main() doesn't use r1-r4?
  - No harm done, but wasted work

```
void foo(){  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    save r1 to stack  
    save r4 to stack  
    bar();  
    restore r1  
    restore r4  
    d = a+d;  
    return();  
}
```

```
void foo(){  
    int a,b,c,d;  
    save r1  
    save r2  
    save r3  
    save r4  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    bar();  
    d = a+d;  
    restore r1  
    restore r2  
    restore r3  
    restore r4  
    return();  
}
```

# Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Caller-saved
  - Only needs saving if value is “**live**” across function call
  - **Live** = contains a useful value: Assign value before function call, use that value after the function call
  - In a leaf function (a function that calls no other function), caller saves can be used without saving/restoring

a, d are live

b, c are NOT live

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

# Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Callee-saved
  - Only needs saving at beginning of function and restoring at end of function
  - Only save/restore it if function overwrites the register

Only use r1-  
r4

No need to  
save other  
registers

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

# Agenda

- Branching far away
- Function calls and the call stack
- Assigning variables to memory locations
- Saving registers
- **Caller/callee example**

# Caller versus Callee

- Which is better??
- Let's look at some examples...
- Simplifying assumptions:
  - A function can be invoked by many different call sites in different functions.
  - Assume no inter-procedural analysis (hard problem)
    - A function has no knowledge about which registers are used in either its caller or callee
    - Assume main() is not invoked by another function
  - Implication
    - Any register allocation optimization is done using function local information



# Caller-saved vs. callee saved – Multiple function case

```
void main() {
    int a,b,c,d;

    c = 5; d = 6;
    a = 2; b = 3;
    foo();
    d = a+b+c+d;

}
```

```
void foo() {
    int e,f;

    e = 2; f = 3;
    bar();
    e = e + f;
```

```
void bar() {
    int g,h,i,j;

    g = 0; h = 1;
    i = 2; j = 3;
    final();
    j = g+h+i;
```

```
void final() {
    int y,z;

    y = 2; z = 3;
    z = y+z;
```

Note: assume main does not have to save any callee registers

# Caller-saved vs. callee saved – Multiple function case

- Questions:
  1. How many registers need to be saved/restored if we use a **caller-save** convention?
  2. How many registers need to be saved/restored if we use a **callee-save** convention?
  3. How many registers need to be saved/restored if we use a mix of **caller-save** and **callee-save**?

# Question 1: Caller-save

```
void main(){
    int a,b,c,d;
    c = 5; d = 6;
    a = 2; b = 3;
    [4 STUR]
    foo();
    [4 LDUR]
    d = a+b+c+d;
}
```

```
void foo(){
    int e,f;
    e = 2; f = 3;
    [2 STUR]
    bar();
    [2 LDUR]
    e = e + f;
}
```

```
void bar(){
    int g,h,i,j;
    g = 0; h = 1;
    i = 2; j = 3;
    [3 STUR]
    final();
    [3 LDUR]
    j = g+h+i;
}
```

```
void final(){
    int y,z;
    y = 2; z = 3;
    z = y+z;
}
```

Total: 9 STUR / 9 LDUR

# Question 2: Callee-save

Poll: How many Id/st pairs are needed?

```
void main() {
    int a,b,c,d;

    c = 5; d = 6;
    a = 2; b = 3;
    foo();
    d = a+b+c+d;

}
```

```
void foo() {
    [2 STUR]
    int e,f;

    e = 2; f = 3;
    bar();
    e = e + f;

}

[2 LDUR]
```

```
void bar() {
    [4 STUR]
    int g,h,i,j;
    g = 0; h = 1;
    i = 2; j = 3;
    final();
    j = g+h+i;

}

[4 LDUR]
```

```
void final() {
    [2 STUR]
    int y,z;

    y = 2; z = 3;
    z = y+z;

}

[2 LDUR]
```

Total: 8 STUR / 8 LDUR

# Is one better?

- **Caller-save** works best when we don't have many live values across function call
- **Callee-save** works best when we don't use many registers overall
- We probably see functions of both kinds across an entire program
- Solution:
  - Use both!
  - E.g. if we have 6 registers, use some (say r0-r2) as **caller-save** and others (say r3-r5) as **callee-save**
  - Now each function can optimize for each situation to reduce saving/restoring
  - Not discussed further in this class

# LEGv8 ABI- Application Binary Interface

- The ABI is an agreement about how to use the various registers
- Not enforced by hardware, just a convention by programmers / compilers
- If you want your code to work with other functions / libraries, **follow these**
- Some register conventions in ARMv8
  - X30 is the **link register** – used to hold return address
  - X28 is **stack pointer** – holds address of top of stack
  - X19-X27 are **callee-saved** – function must save these before writing to them
  - X0-15 are **caller-saved** –function must save live values before call
  - X0-X7 used for **arguments** (memory used if more space is needed)
  - X0 used for **return value**

# Next Time

- Finish Up Function Calls
- Talks about linking – the final puzzle piece of software

